

## Relational Algebra Overview

- *Relational Algebra (RA)*: Formal language for querying/manipulating relational databases, underpinning SQL.

- Purpose: Translates high-level queries into executable operations; optimizes execution.

### Basic Operations

- *Union (U)*: Combines tuples, removes duplicates. Formula:  $(R1 \cup R2)$

Example: ActiveEmployees  $\cup$

RetiredEmployees  $\rightarrow$  unique employees.

⚠ Same schema required.

- *Set Difference (-)*: Tuples in  $(R1)$  not in  $(R2)$ .

Formula:  $(R1 - R2)$

Example: AllEmployees -

RetiredEmployees  $\rightarrow$  current employees.

⚠ Same schema required.

- *Selection ( $\sigma$ )*: Filters tuples by condition ( $c$ ).

Formula:  $(\sigma_c(R))$

Example:  $\sigma_{\text{Salary} > 40000}(\text{Employee}) \rightarrow$

employees earning > \$40K.

Efficiency:  $O(N)$  scan;  $O(\log N)$  with index.

- *Projection ( $\pi$ )*: Selects columns, removes duplicates.

Formula:  $(\pi_{\{A1, \dots, An\}}(R))$

Example:  $\pi_{\{SSN, Name\}}(\text{Employee}) \rightarrow$

unique SSN, Name pairs.

Efficiency:  $O(N)$  hashing;  $O(N \log N)$  sorting.

- *Cartesian Product ( $\times$ )*: Combines all tuples from  $(R1)$  and  $(R2)$ .

Formula:  $(R1 \times R2)$

Example: Employee  $\times$  Dependents  $\rightarrow$  all pairs.

Efficiency:  $O(M \times N)$  I/Os.

- *Renaming ( $\rho$ )*: Changes attribute names.

Formula:  $(\rho_{\{B1, \dots, Bn\}}(R))$

Example:  $\rho_{\text{LastName}}$ ,

SocSecNo(Employee)  $\rightarrow$  renamed attributes.

Efficiency: No I/O.

### Derived Operations

- *Intersection ( $\cap$ )*: Common tuples in both relations.

**Formula:**  $(R1 \cap R2 = R1 - (R1 - R2))$

Example: UnionizedEmployees  $\cap$

RetiredEmployees  $\rightarrow$  common employees.

Efficiency:  $M + N$  reads +  $P$  writes.

Implementation:

1. Sort Merge intersection

2. Hash Base intersection (Cost you'll find in the actual sort method)

- *Theta Join ( $\bowtie_{\theta}$ )*: Combines tuples where condition ( $\theta$ ) is true.

**Formula:**  $(R1 \bowtie_{\text{theta}} R2 = \sigma_{\theta}(R1 \text{ 'times' } R2))$

Example: Sells  $\bowtie_{\{\text{Sells.bar} = \text{Bars.name}\}}$

Bars  $\rightarrow$  matching bar names.

Efficiency:  $O(M \times N)$  I/Os.

- *Natural Join ( $\bowtie$ )*: Joins on common attributes.

**Formula:**  $(R1 \bowtie R2)$

Example: Employee  $\bowtie$  Dependents  $\rightarrow$  matches on shared attributes.

Efficiency:  $M + N$  reads +  $P$  writes (hash join).

- *Equi-Join*: Theta join with equality condition.

**Formula:**  $(R1 \bowtie_{\{A=B\}} R2)$

Example:  $R1 \bowtie_{\{A=B\}} R2 \rightarrow$  joins on  $(A = B)$

).

Efficiency:  $O(M + N)$  hash join.

### Join Overview

*Join*: Combines tuples from relations  $(R)$  and  $(S)$  based on a condition (e.g.,  $(R.X = S.Y)$ ).

**Purpose:** Links related data across tables (e.g., employees with departments).

**Why Output Cost Ignored:** Identical output across methods; varies by tuple size; focus on operational I/Os.

#### Theta Join ( $\bowtie_{\theta}$ )

*Theta Join*: Combines tuples where condition ( $\theta$ ) (e.g.,  $(R.X > S.Y)$ ) is true.

**Formula:**  $(R1 \bowtie_{\theta} R2 = \sigma_{\theta}(R1 * R2))$

**Example:** Sells  $\bowtie_{\{\text{Sells.bar} = \text{Bars.name}\}}$

Bars  $\rightarrow$  matching bar names.

**Implementation:**

Compute Cartesian Product  $(R1 * R2)$ .

Apply selection  $(\sigma_{\theta})$  to filter tuples.

**Estimated Cost:**

I/Os:  $(M + \lceil M / B \rceil * N)$  (Cartesian Product cost).

Computational:  $O(M * N)$  for condition checks.

#### Natural Join ( $\bowtie$ )

*Natural Join*: Joins on common attributes with equal values.

**Formula:**  $(R1 \bowtie R2)$

**Example:** Employee  $\bowtie$  Dependents  $\rightarrow$  matches on shared attributes (e.g., SSN).

**Implementation:**

Identify common attributes.

Perform Equi-Join on these attributes; output combined tuples with single copy of common attributes.

**Estimated Cost:**

I/Os:  $(M + N)$  (hash join);  $(M + \lceil M / B \rceil * N)$  (nested loops).

Computational:  $O(M + N)$  (hash join);  $O(M * N)$  (nested loops).

⚠ Same schema for common attributes required.

#### Equi-Join

*Equi-Join*: Theta Join with equality condition (e.g.,  $R.A = S.B$ )).

**Formula:**  $(R1 \bowtie_{\{A=B\}} R2)$

**Example:**  $R1 \bowtie_{\{A=B\}} R2 \rightarrow$  joins on  $(A = B)$ .

**Implementation:** Similar to Theta Join, but optimize for equality (e.g., use hash table).

**Estimated Cost:**

I/Os:  $(M + N)$  (hash join);  $(M + \lceil M / B \rceil * N)$  (nested loops).

Computational:  $O(M + N)$  (hash join).

#### Database Join Algorithms

##### Simple Nested loop Join

For each tuple in  $R$  we scan through each tuple in  $S$ .

**Cost:**  $M + (Pr * M * N)$ , where  $Pr$  is the number of tuples per page in  $R$ .

##### Page Oriented Nested Loop Join

For each page in  $R$  we scan through each page in  $S$

**Cost:**  $M + (M * N)$

### Block Nested Loop Join

*Block Nested Loop Join*: Processes outer relation in blocks to reduce I/Os.

**Formula (Cost):**  $(R + \lceil R / (B - 2) \rceil * S)$

**Example:**  $(R=10,000, S=8,000, B=2,001)$ , cost 58,000 I/Os (smaller as outer).

**Implementation:**

Allocate  $(B-2)$  buffers for outer relation block, 1 for inner, 1 for output.

For each block of  $(R)$ , build hash table; scan  $(S)$ , probe for matches.

**Estimated Cost:**

I/Os:  $(R + \lceil R / (B - 2) \rceil * S)$

Computational:  $O(R \text{ 'times' } S)$  per block, reduced by hashing.

**Optimization:** Choose smaller relation as outer to minimize  $(\lceil R / (B - 2) \rceil)$

### Sort-Merge Join

*Sort-Merge Join*: Sorts both relations, merges matching tuples.

**Example:**  $(R(\text{sid: 28, 31}), S(\text{sid: 28, 31}))$ , outputs pairs like  $((28, 28))$ .

**Implementation:**

Sort  $(R)$  and  $(S)$  using external sorting.

Merge: Scan both, output matches when  $(R.X = S.Y)$ ; handle duplicates efficiently.

**Estimated Cost:**

Sorting:  $(4M + 4N)$  (2 passes each).

Merging:  $(M + N)$ ; up to  $(M * N)$  with many duplicates.

Total:  $(5M + 5N)$  (typical);  $(4M + 4N + M * N)$  (worst).

**Key Point:** Efficient if pre-sorted; sensitive to duplicates.

⚠ Common mistake: Using in-memory sort cost  $((M \log M))$  instead of external  $((4M))$ .

### Hash Join

*Hash Join*: Partitions relations into buckets, joins matching buckets.

**Formula (Cost):**  $(3M + 3N)$

**Example:**  $(R(\text{sid: 1, 4, 7}))$ , hashed by  $\text{sid mod } 3$ , joins with  $(S)$ .

**Implementation:**

Partition  $(R)$  and  $(S)$  into buckets using hash function ( $h$ ).

For each bucket pair, build hash table for  $(R)$ -bucket, probe with  $(S)$ -bucket.

**Estimated Cost:**

Partitioning:  $(2M + 2N)$  (read/write each relation).

Probing:  $(M + N)$  (read buckets).

Total:  $(3M + 3N)$ .

**Assumption:**  $(R)$ -partitions fit in memory  $((M / (B-1) \leq B-2))$ .

⚠ Common mistake: Applying to non-equi joins (requires equality).

### Index Nested Loop Join

*Index Nested Loop Join*: Uses index on inner relation for efficient lookups.

**Formula (Cost):**  $(M + (M * Pr * \{\text{probe cost}\}))$

**Example:**  $(R(\text{sid: 28}))$ , probes  $(S)$ -index for  $(\text{sid}=28)$ .

**Implementation:**

Scan  $(R)$  sequentially.

For each  $(R)$ -tuple, probe  $(S)$ -index (e.g.,  $B+$  tree) for matches.

**Estimated Cost:**

Read  $(R)$ :  $(M)$  I/Os.

Probes:  $(M * Pr * (\text{Cost for finding Matching tuples}))$  I/Os ( $Pr$ : tuples/page, Cost

for finding matching tuples: 1.2 for Hash Index, 2-4 for  $B+$  tree index).

Index I/Os often cached.

**Assumption:** Index (e.g.,  $B+$  tree) exists on  $(S)$ .

⚠ Common mistake: Assuming efficiency without index.

### Query Building

- *RA Expressions*: Combine operations into sequences.

Example:  $\pi_{\text{name}}(\sigma_{\text{age} > 30}(\text{employee})) \rightarrow$  names of employees > 30.

*Expression Trees*: Visualize queries.

Example:

$\pi_{\text{name}}$

$\downarrow$

$\sigma_{\text{age} > 30}$

$\downarrow$

employee

Leaves: Relations; Root: Result.

Example Query: Bars on Maple St. or selling

Bud < \$3:

**Formula:**  $(\pi_{\text{name}}(\sigma_{\text{addr} = \text{"Maple St."}}(\text{Bars})) \cup \pi_{\text{bar}}(\sigma_{\text{beer} = \text{"Bud" AND price} < 3}(\text{Sells})))$

### Set vs. Bag Semantics

- *Set Semantics*: No duplicates; RA default.

Example:  $\pi_{\text{name}}(\text{employee}) \rightarrow$  one "John". Drawback: Duplicate removal costly.

- *Bag Semantics*: Duplicates allowed; SQL default.

Union: Adds occurrences (e.g.,  $\{a,b,b\} \cup \{b,c\} = \{a,b,b,b,c\}$ ).

Difference: Subtracts occurrences

(e.g.,  $\{x,x,x,y,y,z\} - \{x,y,z,z\} = \{x,x,y\}$ ).

Advantage: Faster, supports streaming.

Example (Bag Difference):

Input:  $(A = \{x,x,x,y,y,z\}, B = \{x,y,z,z\})$

Output:  $(\{x,x,y\})$  using hash map.

### Two-Way External Merge Sort

- *Two-Way Merge Sort*: Simplest external sorting method using 3 buffer pages (2 input, 1 output).

- **Process:**

**Pass 0:** Read 1 page, sort in memory, write as sorted *run*. Repeat for all pages.

**Merge Passes:** Merge 2 runs at a time into a larger sorted run until 1 run remains.

- **Formulas:**

Number of Passes:  $\lceil \log_2 N \rceil + 1$ . ( $N$ : pages)

Total I/O cost:  $2N * (\lceil \log_2 N \rceil + 1)$

- **Example:**

$N=4$ : Pages  $[3,4], [6,2], [9,4], [8,7] \rightarrow$  Runs

$[3,4], [2,6], [4,9], [7,8]$ .

Pass 1: Merge to  $[2,3,4,6], [4,7,8,9]$ .

Pass 2: Merge to  $[2,3,4,6,7,8,9]$ .

Cost: 3 passes,

$2 \times 4 \times 3 = 24$

- **Key Point:** Simple but requires more passes, increasing I/O costs.

### General External Merge Sort

*General Merge Sort*: Uses  $B$  buffer pages to create larger runs and merge more runs per pass.

**Process:**

- **Pass 0:** Read  $B$  pages, sort in memory, write as a *run*. Creates  $\lceil N/B \rceil$  runs.

- **Merge Passes:** Merge  $B-1$  runs at a time using  $B-1$  input buffers and 1 output buffer.

- **Formulas:**

> Number of passes:  $1 + \lceil \log_{B-1} \lceil N/B \rceil \rceil$

> Total I/O cost:  $2N * (1 + \lceil \log_{B-1} \lceil N/B \rceil \rceil)$

### Typical External Merge Sort

### QUERY OPTIMIZATION

**What & Why?**

• What: turn SQL/RA into an efficient

physical plan

• Why: same logical query can vary hugely in

I/O & CPU cost

### Logical vs. Physical Plans

• Logical: tree of  $\sigma, \pi, \bowtie, \cup$ , etc.

• Physical: choose an algorithm per operator (nested-loop, hash-join, sort-merge, etc.) +

buffer/pipeline strategy

### Pipelining vs. Materialize

• Pipeline: feed tuples directly to next

operator—no temp writes

• Materialize: write intermediate to disk, then

read back

### “Good” Plans & Plan Space

• Join trees:

– Left-deep: outer input is a base table or

prior join (easy pipelining)

– Right-deep / bushy: more parallelism,

complex buffers

• 3 tables  $\rightarrow$  6 left-deep orders

• 3 join algls  $\rightarrow 6 \times 5^2 = 150$  plans (before

pruning)

### Push-down Rules

1. Selection push-down: apply  $\sigma$  as early as possible

2. Projection push-down: drop unused cols at leaves

3. Push  $\sigma$  into indexed NL: e.g.

$A \bowtie_{\sigma(B.\text{price} > 5)}(B)$  uses  $B.\text{id}$  index

### Cost Estimation (I/O)

• Scan/ $\sigma/\pi$ :  $B(R)$

• Simple NL:  $B(R) + T(R) \times B(S)$

• Block NL:  $B(R) + [B(R)/(B-2)] \times B(S)$

• Index NL:  $B(\text{outer}) + T(\text{outer}) \times (\text{index}$

lookup)

• Sort-merge join (2-pass if  $B^2 > \text{pages}$ ):

- Sort R:  $2 \times B(R) \cdot [1 + \lceil \log_{B-1} (B(R)/B) \rceil]$

- Sort S: same

- Merge:  $B(R) + B(S)$

• Hash-join  $(M < (B-2)^2)$ : partition  $2M + 2N$ ,

build/probe  $M + N \Rightarrow 3M + 3N$

### 2. TRANSACTION MANAGEMENT

#### What is a Transaction?

A sequence of SQL statements executed atomically (all or none).

#### Why Need?

Prevent inconsistency from crashes or concurrent updates (e.g. \$20 transfer between accounts).

### SQL Syntax

BEGIN TRANSACTION;

...SQL statements...  
COMMIT; -- or ROLLBACK

### ACID Properties & Who Enforces

- Atomicity - All-or-nothing execution - enforced by DBMS
- Consistency - DB moves from one valid state to another - enforced by programmer (App logic + constraints)
- Isolation - Concurrent Ts appear serial - enforced by DBMS
- Durability - Once committed, changes survive crashes - enforced by DBMS

### Locks & Concurrency Control

- Granularity: database > table > row > cell
- Pessimistic: lock before access
- Optimistic: validate at commit

### Crash Recovery Intro

Primitive ops per T:  
INPUT(X) : read page X into buffer  
READ(X,t) : buffer X → local t  
WRITE(X,t) : local t → buffer X  
OUTPUT(X) : flush buffer X to disk

### 3. RECOVERY & CHECKPOINTING

Failures: system crashes (power loss) ⇒ memory lost, disk partial writes.  
Assumptions: WAL ⇒ log records forced before data pages.

#### A) Undo-Logging

Log records:  
<START T>  
<T, X, oldValue>  
<COMMIT T>  
WAL rules:  
1. Write <T,X,old> before OUTPUT(X)  
2. Write <COMMIT T> after all OUTPUTS of T  
Fuzzy checkpoint:  
<START\_CKPT.Active={T1...Tk}>  
...flush pages of already-committed Ts...  
<END\_CKPT>  
Undo recovery:  
1. Backward scan to last <END\_CKPT>, then to its <START\_CKPT>; let C=Active.  
2. Forward scan from <START\_CKPT> to crash:

- Mark committed Ts seen
- Any TεC w/o commit, and any new T' w/o commit ⇒ ToUndo
- 3. Backward scan crash→<START\_CKPT>:
  - For each <T,X,old> with TεToUndo: X:=old
  - On <START T>: remove T from ToUndo

#### B) Redo-Logging

Log records:  
<START T>  
<T, X, newValue>  
<COMMIT T> -- may precede data writes  
WAL rule:  
Before OUTPUT(X), <T,X,new> & <COMMIT T> must be on disk  
Fuzzy checkpoint: same markers as undo  
Redo recovery:  
1. Backward scan to last <END\_CKPT>, then to its <START\_CKPT>  
2. Forward scan <START\_CKPT>→crash: collect committed Ts ⇒ ToRedo  
3. Forward scan again: for each <T,X,new> with TεToRedo: X:=new

#### C) Combined (ARIES-style) [optional]

- Analysis: identify Dirty pages & active Ts
- Redo: replay updates from earliest dirty LSN

3. Undo: rollback Ts still active at crash

### 1. Relational Algebra

(a) Two Derived Operations of Relational Algebra  
**Answer:** Intersection and Join  
(b) Expressing Intersection in Terms of Basic Operations  
**Answer:**  $R - (R - S)$

### 2. Implementation of Operators

(a) Union Operator Using Sort-Merge  
**Answer:** Sort-merge union with cost  $5M + 5N$ .  
Explanation:  
 $M, N < B(B-1)$ , so sorting completes in two passes.  
Sorting (External Sort):  
- Pass 0: Divide R and S into sorted runs (read/write each page once →  $2M + 2N$ ).  
- Pass 1: Merge runs into final sorted files (read/write again →  $2M + 2N$ ).  
- Total sorting cost:  $4M + 4N$ .  
Merging:  
- Merge sorted R and S, skipping duplicates (read each page once →  $M + N$ ).  
Total Cost:  $4M + 4N$  (sorting) +  $M + N$  (merging) =  $5M + 5N$ .

(b) Two Hash Functions in Hash Join  
**Answer:** One hash function partitions relations into B-1 buckets; a different one builds in-memory hash tables for uniform distribution.  
Explanation:  
First Hash Function: Used to divide R and S into B-1 buckets on disk. Ensures tuples with matching keys (same join attribute) go to the same bucket.  
Second Hash Function: Used within each bucket to build an in-memory hash table (tuned to RAM size). Ensures uniform distribution of tuples for fast lookups during probing.  
Why Different: The first function optimizes for disk partitioning; the second for in-memory efficiency. Using the same function could lead to poor in-memory distribution, slowing lookups.

(c) Cost of Hash Join

**Answer:** Ideal cost is  $3M + 3N$ .  
Steps (Ideal Case):  
Partitioning: Read R and S, hash into B-1 buckets, write back (read + write →  $2M + 2N$ ).  
Build and Probe: Read each bucket of R and S into memory (fits due to  $M < (B-2)(B-2)$ ), build hash table, and probe (read →  $M + N$ ).  
Total Cost:  $2M + 2N$  (partition) +  $M + N$  (build/probe) =  $3M + 3N$ .  
Each page is read twice (partition, probe) and written once (partition), and buckets fit in memory.

(d) Cost of Sort-Merge Join for R and S

**Answer:** Total cost is  $150 I/Os$ .  
Explanation:  
Relations:  
- R (Enroll: sid, cid, grade) = 20 pages, S (Students: sid, sname) = 10 pages.  
- Buffer = 6 pages, each page holds 100 tuples, each student in S enrolls in ≤3 courses.  
Sort-Merge Join Steps:  
> Sorting:  
- R (20 pages) and S (10 pages) <  $B^2$  ( $6^2 = 36$ ), so sorting completes in 2 passes.  
- Cost per relation:  $2 \text{ passes} \times (\text{read} + \text{write}) = 4 \times |\text{pages}|$ .  
- R:  $4 \times 20 = 80 I/Os$ , S:  $4 \times 10 = 40 I/Os$ .

- Total sorting cost:  $80 + 40 = 120 I/Os$ .  
> Merging:  
- Read sorted R and S, join on sid.  
- Each student (tuple in S) matches ≤3 tuples in R (≤3 courses), fitting in one page.  
- Cost: Read R (20 pages) + S (10 pages) = 30 I/Os.  
Total Cost:  $120$  (sorting) +  $30$  (merging) =  $150 I/Os$ .  
Sorting is efficient (2 passes), and merging is straightforward since matches are small (≤3 per student).

### 3. Query Optimization

(a) 12 Physical Query Plans for SQL Query  
**Answer:** 12 plans: 6 join orders ( $A \rightarrow B \rightarrow C$ ,  $A \rightarrow C \rightarrow B$ ,  $B \rightarrow A \rightarrow C$ ,  $B \rightarrow C \rightarrow A$ ,  $C \rightarrow A \rightarrow B$ ,  $C \rightarrow B \rightarrow A$ )  $\times$  2 join algorithms (from simple nested-loop, block nested-loop, index nested-loop, hash join, sort-merge join). Push selections ( $B.\text{price} > 5$ ,  $C.\text{age} > 30$ ) and projections early.  
**Explanation:**  
**Query:** Joins A, B, C on A.id = B.id and B.id = C.id, filters B.price > 5, C.age > 30, projects A.name.  
**Plans:**  
**Join Orders:** 6 permutations ( $3! = 6$ ).  
**Join Algorithms:** Pick 2 of 5 (simple/block/index nested-loop, hash, sort-merge).  
**Total:**  $6 \times 2 = 12$  plans.

**Optimizations:**  
Apply selections ( $B.\text{price} > 5$ ,  $C.\text{age} > 30$ ) before joins.  
Project only required columns early.  
Use join conditions (A.id = B.id, B.id = C.id, or A.id = C.id).  
**Purpose:** Vary orders/algorithms to minimize I/O and CPU costs.

(b) Pushing Selection Below Indexed Join  
**Answer:** Yes, push B.price > 5 below the join, as it doesn't affect the B.id index.  
**Explanation:**  
**Context:**  $A \bowtie B$  with indexed nested-loop join using B.id index.  
**Selection:** B.price > 5 filters on price, not id.  
**Reason:**  
- Selection is independent of B.id, preserving index usability.  
- Equivalent:  $\sigma_{B.\text{price} > 5}(A \bowtie B) = A \bowtie \sigma_{B.\text{price} > 5}(B)$ .  
**Benefit:** Reduces B's tuples before join, lowering cost.

### 4. Transaction Management

(a) What is a Transaction in Relational Databases?  
**Answer:** A transaction is a sequence of SQL statements executed as a single atomic unit—either all succeed or none are applied.  
**Explanation:**  
**Definition:** A transaction groups SQL operations (e.g., INSERT, UPDATE) to act as one unit.  
**Atomicity:** Ensures all statements complete successfully, or no changes are made (e.g., if one fails, all are undone).  
**Consistency:** Maintains database validity—starting consistent, it ends consistent.  
**Purpose:** Prevents partial execution, which could leave the database in an inconsistent state.  
**Execution:** Transactions run sequentially in the application.

(b) Transaction V Interfering with Transaction T  
**Answer:** Transaction V checks the account balance (SELECT checking + savings). Without ACID, V may read an inconsistent balance if it runs between T's subtraction from checking and addition to savings.  
**Explanation:**  
**Transaction T:** Moves \$20 from checking to savings:  
- Subtract \$20 from checking.  
- Add \$20 to savings.  
**Transaction V:** Queries total balance:  
**SQL:** SELECT (checking + savings) AS total FROM Accounts WHERE acct\_id = 123;  
**Interference:**  
- Without ACID (Isolation), V could run after T subtracts \$20 from checking but before adding to savings.  
- V would see a lower balance (missing the \$20 not yet added to savings), causing inconsistency.  
**Why ACID Matters:** Isolation ensures V sees either the state before or after T, not an intermediate state.

(c) How ACID Requirements Are Ensured  
**Answer:**  
**Atomicity:** DBMS uses transaction logs to undo partial changes if a transaction fails.  
**Consistency:** Integrity constraints and rules enforce valid state transitions.  
**Isolation:** Concurrency controls (e.g., locking) ensure transactions appear to run serially.  
**Durability:** Write-ahead logging ensures committed changes are saved to stable storage.

### 5. Recovery

(a) Non-Quiescent Checkpointing  
**Answer:** Non-quiescent checkpointing logs active transactions (<START CKPT (T1,...,Tk)>) without pausing the database, flushes committed dirty pages, and logs <END CKPT> when active transactions complete. The database does not freeze.  
**Explanation:**  
**Concept:** Records a checkpoint without stopping new transactions.  
**Process:**  
- Log <START CKPT (T1,...,Tk)> with active transactions.  
Continue normal operations; flush committed dirty pages in the background.  
Log <END CKPT> after T1,...,Tk commit or rollback.  
**Recovery:** After a crash, start from last <END CKPT>, process only listed transactions and later ones.  
**No Freeze:** Only log records are written atomically, allowing ongoing transactions.  
**Purpose:** Reduces recovery time while maintaining consistency without halting the system.

(b) Undo Logging: Transactions to Undo After Crash  
**Answer:** Undo active transactions (with <START T> but no <COMMIT T>) because their partial changes may be on disk, risking inconsistency.  
**Explanation:**  
- **Transactions to Undo:** Active transactions at crash: Have <START T> but no <COMMIT T>. Their changes may be on disk (undo logging allows early writes), leaving an inconsistent state.

- **Recovery Process:**  
Scan log backward.  
Restore old page values for these transactions' updates (rollback).

**Why Not Others:**  
Committed transactions (<COMMIT T> logged) have all changes safely on disk.

**Purpose:** Ensures atomicity and consistency by removing partial, uncommitted changes.

6. Recovery  
(a) Transactions Flushed After <start ckpt t5,t6>  
**Answer:** T1, T4

**Explanation:**  
**Context:** After logging <start ckpt t5,t6>, the system flushes changes by committed transactions.  
**Transactions:**  
T1: Committed (<commit t1>) before checkpoint.  
T4: Committed (<commit t4>) before checkpoint.  
T5, T6: Active during checkpoint (listed in <start ckpt t5,t6>), not yet committed.  
T7: Starts after checkpoint, not relevant yet.

**Why T1, T4:**  
Only committed transactions' changes (dirty pages) are flushed during checkpointing to ensure durability.  
T5, T6 are active, so their changes are not flushed yet.  
**Purpose:** Flushing T1, T4 ensures their committed changes are on disk for recovery.

(b) Recovery Actions Using Redo Log  
**Answer:** Redo T7's update (<t7, z, 4>). Final values: x = 5, y = 3, z = 4.  
**Explanation:**  
**Recovery Steps:**  
**Locate Checkpoint:** Find latest <end ckpt> (indicates completed checkpoint with T5, T6 active).

**Identify Transactions to Redo:**  
- Transactions in <start ckpt t5,t6> that committed: None (T5, T6 have no <commit>).  
- Transactions started after <start ckpt> and committed: T7 (<commit t7>).  
**Redo Updates:**  
- Apply T7's log entry: <t7, z, 4> → set z = 4.  
- Flush changes to disk.  
- **Ignore Uncommitted:** T5, T6 have no <commit>, so their changes (e.g., <t5, x, 2>) were not flushed to disk.  
**Final Values:**  
x = 5 (from T1's <t1, x, 5>, committed and flushed).  
y = 3 (from T4's <t4, y, 3>, committed and flushed).  
z = 4 (from T7's <t7, z, 4>, redone during recovery).  
**Why it works:** Redo logging only applies committed changes; uncommitted ones are ignored as they weren't written to disk.